

## Question 3: Binary Search Tree vs Hash Table insertion times

- Hypothesis:


In the C++ standard library, `std::multiset` can be used to easily store values in a balanced binary tree. Likewise, `std::unordered_multiset` can be used to store values in a hash table. In theory, a hash table should be faster for insertions and deletions. We predict this to be true, and for hash table insertions to be roughly 50% faster in large values of  $n$ .

- Methods:

We inserted  $n$  values into both a BST and a hash table at varying values of  $n$ , starting at 10. We then timed how long it took to run  $n$  amount of inserts, with  $n$  being multiplied by 10 every time.

- Results:

We ran various experiments with different source code, and ended up concluding that at  $n = 1000000$ , the hash table was 25% faster (10.799 seconds to insert 1000000 items compared to 8.635 seconds).



- Discussion:

This was about what we had predicted. Something that was surprising was that at  $n=1,000$  and  $n=10,000$ , the insertion time for the hash table is slightly slower than the BST. This is likely due to  $n$  being too small to get an accurate measurement.

- Conclusions:

Under the conditions tested, insertion time for hash tables is consistently faster than insertion time for binary search trees.