

Question 4: Multiset (BST) vs Vectors

- Hypothesis:

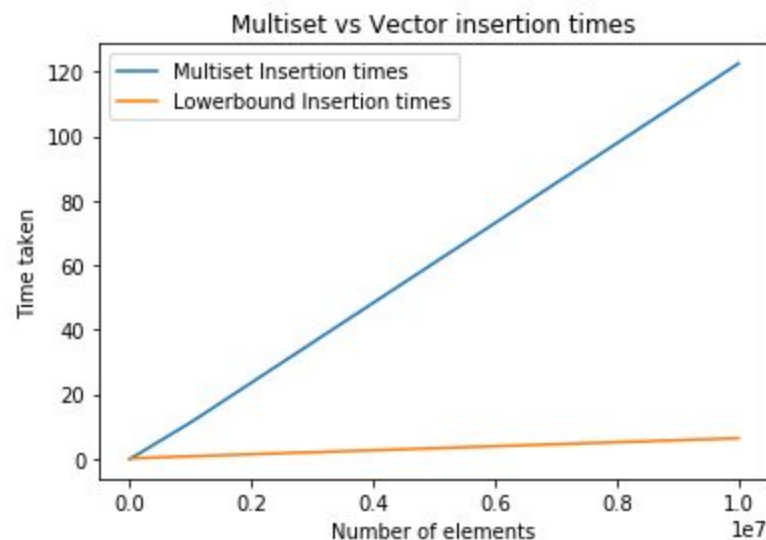
In C++, `std::vector` is typically considered a poor choice for a sorted container because for every insertion, all subsequent values must be shifted to make room for it. The net effect should be linear-time insertions. We predict that Multiset insertion times will be marginally slower than insertion times for a vector using binary search.

- Methods:

We inserted n values into BST and found the lower bound position for sorted vector and then insert it into that position, at varying values of n , starting at 10. We then timed how long it took to run n amount of inserts, with n being multiplied by 10 every time.

- Results:

We ran various experiments with different source code, and ended up concluding that all



values of n , inserting into the vector was roughly 99% faster than inserting into a multiset.

n	Multiset insertio n time (sec)	Vector insertion time (sec)
10	0	0
100	0.001	0
1000	0.011	0
10000	0.106	0.001
100000	1.2	0.002

1000000	10.47	0.011
10000000	129.763	0.198

- Discussion:
This is not at all what we had predicted. We were very far off as we thought that inserting into a multiset was only going to be marginally slower than inserting into a vector. The time to insert into a vector never got over 2/10ths of a second, while insertion time for the multiset was exceeding 100 seconds.
- Conclusions:
Under the conditions tested, insertion time for vector using binary search is always and significantly faster than inserting into a multiset.